

Importing Libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

Reading dataset

```
dataset = pd.read_csv("Car_Price_dataset.csv")
dataset
```

	car_ID	symboling	CarName	fueltype	
aspiration \					
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std
..
200	201	-1	volvo 145e (sw)	gas	std
201	202	-1	volvo 144ea	gas	turbo
202	203	-1	volvo 244dl	gas	std
203	204	-1	volvo 246	diesel	turbo
204	205	-1	volvo 264gl	gas	turbo
doornumber	carbody	drivewheel	engine	location	wheelbase ...
\					
0	two	convertible	rwd	front	88.6 ...
1	two	convertible	rwd	front	88.6 ...
2	two	hatchback	rwd	front	94.5 ...

3	four	sedan	fwd	front	99.8	...
4	four	sedan	4wd	front	99.4	...
..
200	four	sedan	rwd	front	109.1	...
201	four	sedan	rwd	front	109.1	...
202	four	sedan	rwd	front	109.1	...
203	four	sedan	rwd	front	109.1	...
204	four	sedan	rwd	front	109.1	...
<div> <div>enginesize</div> <div>fuelsystem</div> <div>boreratio</div> <div>stroke</div> <div>compressionratio</div> <div>horsepower \</div> </div>						
0	130	mpfi	3.47	2.68	9.0	
111						
1	130	mpfi	3.47	2.68	9.0	
111						
2	152	mpfi	2.68	3.47	9.0	
154						
3	109	mpfi	3.19	3.40	10.0	
102						
4	136	mpfi	3.19	3.40	8.0	
115						
..	
...						
200	141	mpfi	3.78	3.15	9.5	
114						
201	141	mpfi	3.78	3.15	8.7	
160						
202	173	mpfi	3.58	2.87	8.8	
134						
203	145	idi	3.01	3.40	23.0	
106						
204	141	mpfi	3.78	3.15	9.5	
114						
<div> <div>peakrpm</div> <div>citympg</div> <div>highwaympg</div> <div>price</div> </div>						
0	5000	21.0	27.0	13495.0		
1	5000	21.0	27.0	16500.0		
2	5000	NaN	26.0	16500.0		
3	5500	24.0	30.0	13950.0		
4	5500	18.0	22.0	17450.0		
..		
200	5400	23.0	28.0	16845.0		

201	5300	19.0	25.0	19045.0
202	5500	18.0	23.0	21485.0
203	4800	26.0	27.0	22470.0
204	5400	19.0	25.0	22625.0

[205 rows x 26 columns]

dataset.dtypes

car_ID	int64
symboling	int64
CarName	object
fueltype	object
aspiration	object
doornumber	object
carbody	object
drivewheel	object
enginelocation	object
wheelbase	float64
carlength	float64
carwidth	float64
carheight	float64
curbweight	float64
enginetype	object
cylindernumber	object
enginesize	int64
fuelsystem	object
boreratio	float64
stroke	float64
compressionratio	float64
horsepower	int64
peakrpm	int64
citympg	float64
highwaympg	float64
price	float64

dtype: object

dataset.isna().sum()

car_ID	0
symboling	0
CarName	0
fueltype	0
aspiration	0
doornumber	0
carbody	0
drivewheel	0
enginelocation	0
wheelbase	5
carlength	0

```

carwidth      3
carheight     3
curbweight    1
enginetype    0
cylindernumber 0
enginesize    0
fuelsystem    0
boreratio     4
stroke        1
compressionratio 2
horsepower    0
peakrpm       0
citympg       3
highwaympg    2
price         0
dtype: int64

```

```
dataset.describe()
```

```

      car_ID  symboling  wheelbase  carlength  carwidth
carheight \
count  205.000000  205.000000  200.000000  205.000000  202.000000
202.000000
mean   103.000000    0.834146   98.816500  174.049268   65.928218
53.731683
std     59.322565    1.245307    6.079414   12.337289    2.152159
2.451801
min      1.000000   -2.000000   86.600000  141.100000   60.300000
47.800000
25%     52.000000    0.000000   94.500000  166.300000   64.125000
52.000000
50%    103.000000    1.000000   97.000000  173.200000   65.500000
54.100000
75%    154.000000    2.000000  102.400000  183.100000   66.900000
55.500000
max    205.000000    3.000000  120.900000  208.100000   72.300000
59.800000

```

```

      curbweight  enginesize  boreratio  stroke
compressionratio \
count  204.000000  205.000000  201.000000  204.000000
203.000000
mean   2554.808824  126.907317    3.330100    3.254412
10.155271
std     521.847984   41.642693    0.272101    0.314039
3.989187
min   1488.000000   61.000000    2.540000    2.070000
7.000000
25%    2145.000000   97.000000    3.150000    3.110000
8.600000

```

50%	2414.000000	120.000000	3.310000	3.290000
9.000000				
75%	2939.250000	141.000000	3.580000	3.410000
9.400000				
max	4066.000000	326.000000	3.940000	4.170000
23.000000				

	horsepower	peakrpm	citympg	highwaympg	price
count	205.000000	205.000000	202.000000	203.000000	205.000000
mean	104.117073	5125.121951	25.277228	30.822660	13276.710571
std	39.544167	476.985643	6.515376	6.880744	7988.852332
min	48.000000	4150.000000	13.000000	16.000000	5118.000000
25%	70.000000	4800.000000	19.000000	25.000000	7788.000000
50%	95.000000	5200.000000	24.000000	30.000000	10295.000000
75%	116.000000	5500.000000	30.000000	35.000000	16503.000000
max	288.000000	6600.000000	49.000000	54.000000	45400.000000

Replacing missing values

```
for column, content in dataset.items():
    if pd.isna(content).sum():
        content.fillna(content.mean(), inplace=True)
```


```
dataset.isna().sum()
```

car_ID	0
symboling	0
CarName	0
fueltype	0
aspiration	0
doornumber	0
carbody	0
drivewheel	0
enginelocation	0
wheelbase	0
carlength	0
carwidth	0
carheight	0
curbweight	0
enginetype	0
cylindernumber	0
enginesize	0
fuelsystem	0
boreratio	0
stroke	0
compressionratio	0
horsepower	0
peakrpm	0

```
citympg      0
highwaympg   0
price        0
dtype: int64

plt.scatter(x=dataset.CarName, y=dataset.price)
#plt.scatter(dataset.CarName, dataset.price)

<matplotlib.collections.PathCollection at 0x16412bc20>
```

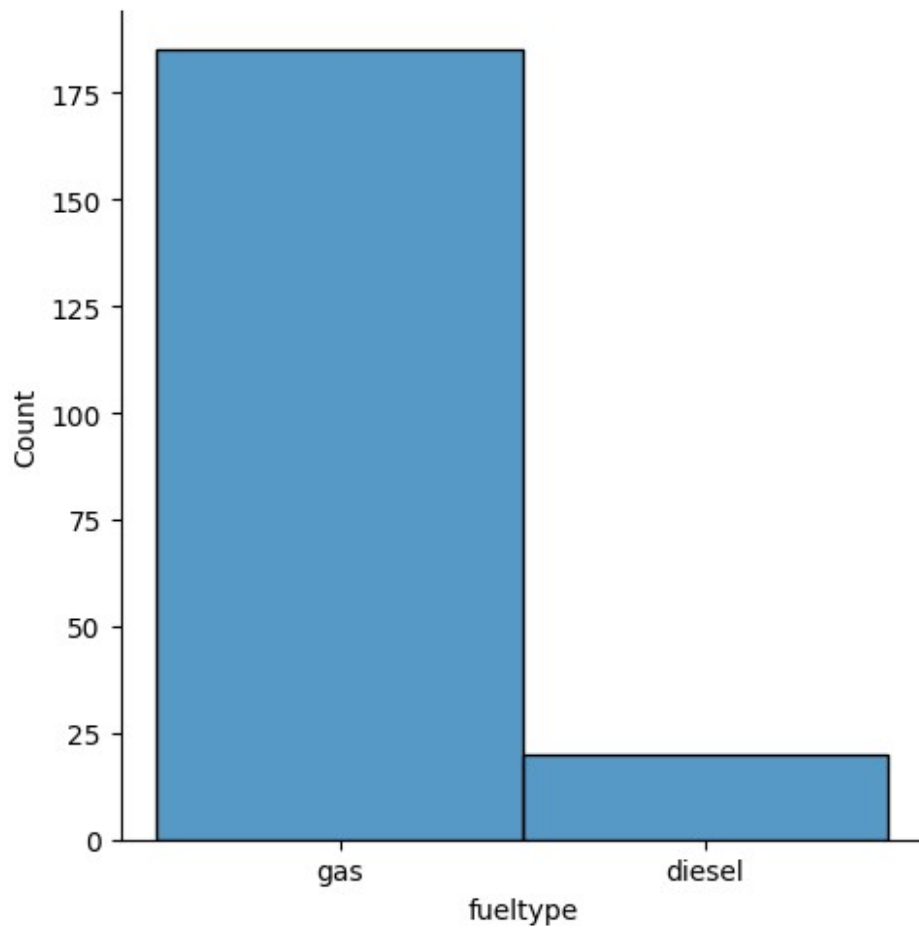


Plotting graphs

1. FuelType
2. DoorNo
3. EngineLocation
4. CarBody

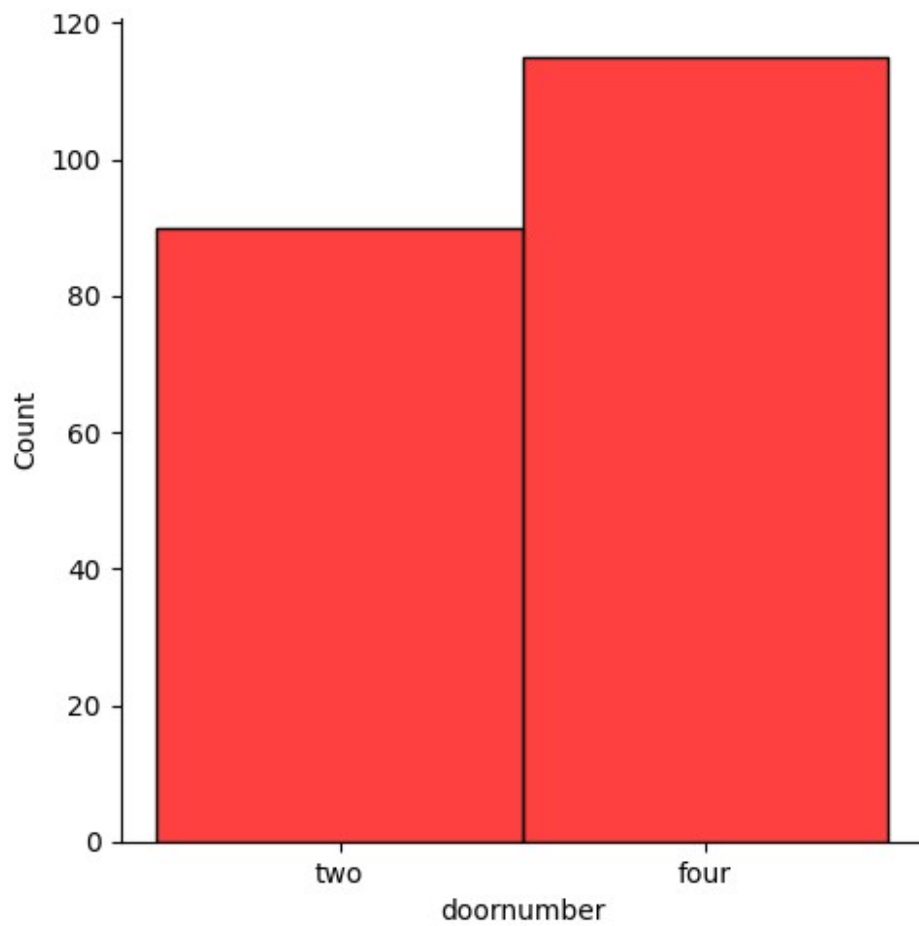
```
sns.displot(dataset, x='fueltype')

<seaborn.axisgrid.FacetGrid at 0x16410b170>
```

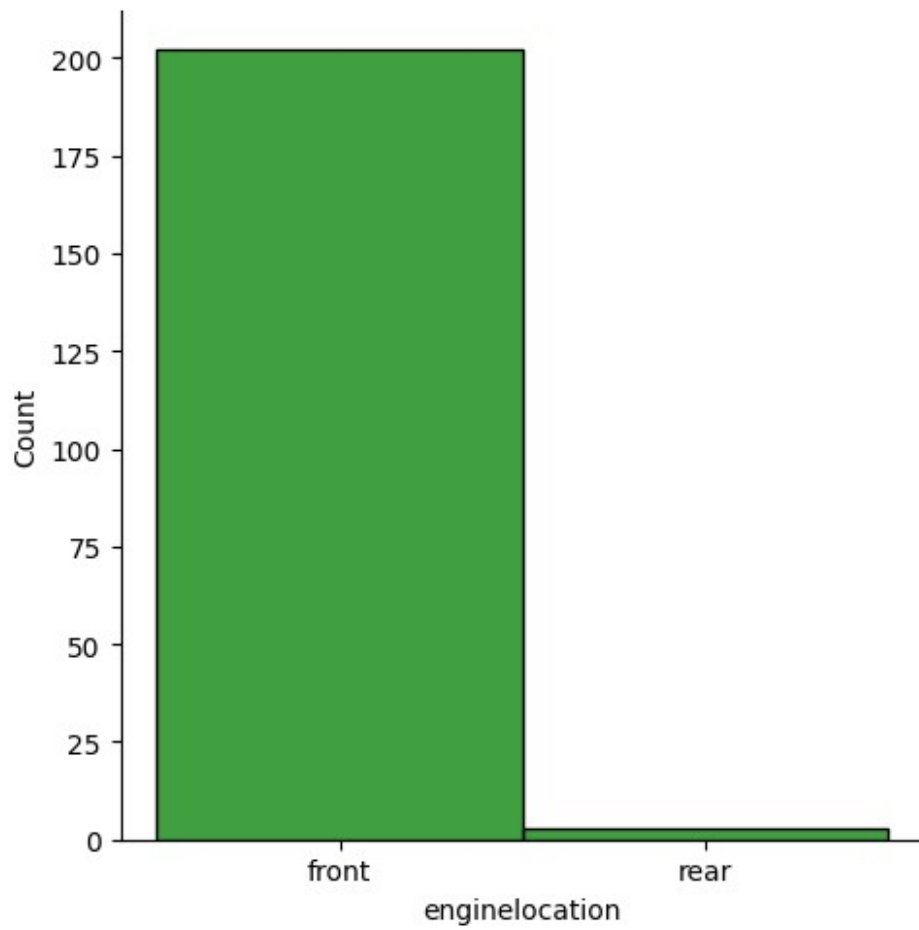


```
sns.displot(dataset, x='doornumber', color='red')
```

```
<seaborn.axisgrid.FacetGrid at 0x1653272f0>
```




```
sns.displot(dataset, x='enginelocation', color='green')  
<seaborn.axisgrid.FacetGrid at 0x165536210>
```

```
sns.displot(dataset, x='carbody', color='hotpink')
```

```
<seaborn.axisgrid.FacetGrid at 0x164684a40>
```



```
data = pd.DataFrame(dataset, columns=['engine_size', 'price'])
data.cov()
```

	engine_size	price
engine_size	1734.113917	2.908082e+05
price	290808.157690	6.382176e+07

Making changes in the dataset

```
dataset.drop("car_ID", axis=1)
```

symboling	CarName	fueltype	aspiration
door number 3	alfa-romero giulia	gas	std
1 two 3	alfa-romero stelvio	gas	std
2 two 1	alfa-romero Quadrifoglio	gas	std

3	2	audi 100 ls	gas	std		
four						
4	2	audi 100ls	gas	std		
four						
..
.						
200	-1	volvo 145e (sw)	gas	std		
four						
201	-1	volvo 144ea	gas	turbo		
four						
202	-1	volvo 244dl	gas	std		
four						
203	-1	volvo 246	diesel	turbo		
four						
204	-1	volvo 264gl	gas	turbo		
four						
carbody drivewheel enginelocation wheelbase carlength ...						
\						
0	convertible	rwd	front	88.6	168.8	...
1	convertible	rwd	front	88.6	168.8	...
2	hatchback	rwd	front	94.5	171.2	...
3	sedan	fwd	front	99.8	176.6	...
4	sedan	4wd	front	99.4	176.6	...
..
200	sedan	rwd	front	109.1	188.8	...
201	sedan	rwd	front	109.1	188.8	...
202	sedan	rwd	front	109.1	188.8	...
203	sedan	rwd	front	109.1	188.8	...
204	sedan	rwd	front	109.1	188.8	...
enginesize fuelsystem boreratio stroke compressionratio						
horsepower \						
0	130	mpfi	3.47	2.68	9.0	
111						
1	130	mpfi	3.47	2.68	9.0	
111						
2	152	mpfi	2.68	3.47	9.0	
154						

3	109	mpfi	3.19	3.40	10.0
102					
4	136	mpfi	3.19	3.40	8.0
115					
..
...					
200	141	mpfi	3.78	3.15	9.5
114					
201	141	mpfi	3.78	3.15	8.7
160					
202	173	mpfi	3.58	2.87	8.8
134					
203	145	idi	3.01	3.40	23.0
106					
204	141	mpfi	3.78	3.15	9.5
114					

	peakrpm	citympg	highwaympg	price
0	5000	21.000000	27.0	13495.0
1	5000	21.000000	27.0	16500.0
2	5000	25.277228	26.0	16500.0
3	5500	24.000000	30.0	13950.0
4	5500	18.000000	22.0	17450.0
..
200	5400	23.000000	28.0	16845.0
201	5300	19.000000	25.0	19045.0
202	5500	18.000000	23.0	21485.0
203	4800	26.000000	27.0	22470.0
204	5400	19.000000	25.0	22625.0

[205 rows x 25 columns]

Mapping different fueltypes wrt price

```
data = pd.DataFrame(dataset, columns=['fueltype', 'price'])
data = dataset.fueltype.eq('gas').mul(1)
data
```

0	1
1	1
2	1
3	1
4	1
..	..
200	1
201	1
202	1
203	0

```
204      1
Name: fueltype, Length: 205, dtype: int64
```

```
dataset
```

	car_ID	symboling	CarName	fueltype	
aspiration \					
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std
..
200	201	-1	volvo 145e (sw)	gas	std
201	202	-1	volvo 144ea	gas	turbo
202	203	-1	volvo 244dl	gas	std
203	204	-1	volvo 246	diesel	turbo
204	205	-1	volvo 264gl	gas	turbo

	doornumber	carbody	drivewheel	engine	location	wheelbase	...
\							
0	two	convertible	rwd		front	88.6	...
1	two	convertible	rwd		front	88.6	...
2	two	hatchback	rwd		front	94.5	...
3	four	sedan	fwd		front	99.8	...
4	four	sedan	4wd		front	99.4	...
..
200	four	sedan	rwd		front	109.1	...
201	four	sedan	rwd		front	109.1	...
202	four	sedan	rwd		front	109.1	...
203	four	sedan	rwd		front	109.1	...

204	four	sedan	rwd	front	109.1	...
-----	------	-------	-----	-------	-------	-----

	engine size	fuel system	bore ratio	stroke	compression ratio
0	130	mpfi	3.47	2.68	9.0
111					
1	130	mpfi	3.47	2.68	9.0
111					
2	152	mpfi	2.68	3.47	9.0
154					
3	109	mpfi	3.19	3.40	10.0
102					
4	136	mpfi	3.19	3.40	8.0
115					
..
...					
200	141	mpfi	3.78	3.15	9.5
114					
201	141	mpfi	3.78	3.15	8.7
160					
202	173	mpfi	3.58	2.87	8.8
134					
203	145	idi	3.01	3.40	23.0
106					
204	141	mpfi	3.78	3.15	9.5
114					

	peak rpm	city mpg	highway mpg	price
0	5000	21.000000	27.0	13495.0
1	5000	21.000000	27.0	16500.0
2	5000	25.277228	26.0	16500.0
3	5500	24.000000	30.0	13950.0
4	5500	18.000000	22.0	17450.0
..
200	5400	23.000000	28.0	16845.0
201	5300	19.000000	25.0	19045.0
202	5500	18.000000	23.0	21485.0
203	4800	26.000000	27.0	22470.0
204	5400	19.000000	25.0	22625.0

[205 rows x 26 columns]

data.value_counts()

fueltype

1 185

0 20

Name: count, dtype: int64

Splitting data in labels and feature

Uni-variate

```
X = pd.DataFrame(dataset, columns=['engine_size'])
y = pd.DataFrame(dataset, columns=['price'])

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25)

reg = LinearRegression()
reg.fit(X_train, y_train)

y_preds = reg.predict(X_test)
```

Evaluating the model

```
from sklearn.metrics import mean_squared_error, r2_score,
mean_absolute_error, accuracy_score


mae = mean_absolute_error(y_test, y_preds)
r2 = r2_score(y_test, y_preds)
mse = mean_squared_error(y_test, y_preds)

mae, r2, mse

(2890.5524675075594, 0.8114416757838671, 17215055.277613536)

plt.scatter(X_train, y_train, c='blue')
plt.plot(X_test, y_preds, c='red')

[<matplotlib.lines.Line2D at 0x166032120>]
```



Predicting on random input

```
y_random = y_test.iloc[12]
x_random = X_test.iloc[12]

y_random_pred = reg.predict(x_random.to_numpy().reshape(-1,1))
print(f"Predicted: {y_random_pred}, Og {y_random}")
```

```
Predicted: [[14474.49963978]], Og price 17710.0
Name: 6, dtype: float64
```

```
/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/base.py:493: UserWarning: X does not have valid
feature names, but LinearRegression was fitted with feature names
warnings.warn(
```

Multi-variate

```
X = pd.DataFrame(dataset, columns=['engine_size', 'bore_ratio',
'stroke'])
y = pd.DataFrame(dataset, columns=['price'])
```



```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25)

multi_reg = LinearRegression()
multi_reg.fit(X_train, y_train)
y_preds = multi_reg.predict(X_test)

mae = mean_absolute_error(y_test, y_preds)
r2 = r2_score(y_test, y_preds)
mse = mean_squared_error(y_test, y_preds)

mae, r2, mse
#(2755.219116212158, 0.7459019410391048, 15174421.15204328)

(3176.2263393213525, 0.7137335022864753, 17049992.555727363)

y_random = y_test.iloc[13]
x_random = pd.DataFrame(X_test).iloc[13]
y_random_pred = reg.predict(x_random.to_numpy().reshape(-1,1))
print(f"Predicted: {y_random_pred[0]}, 0g {y_random}")

Predicted: [20053.3502221], 0g price      15998.0
Name: 179, dtype: float64

/Library/Frameworks/Python.framework/Versions/3.12/lib/python3.12/
site-packages/sklearn/base.py:493: UserWarning: X does not have valid
feature names, but LinearRegression was fitted with feature names
warnings.warn(

```

OneHotEncoder

```

dataset.dtypes

car_ID          int64
symboling       int64
CarName         object
fueltype       object
aspiration      object
doornumber      object
carbody         object
drivewheel      object
enginelocation  object
wheelbase      float64
carlength      float64
carwidth       float64
carheight      float64
curbweight     float64
enginetype     object
cylindernumber  object

```

```

enginesize      int64
fuelsystem      object
boreratio       float64
stroke          float64
compressionratio float64
horsepower      int64
peakrpm         int64
citympg         float64
highwaympg      float64
price           float64
dtype: object

```

```

dataset = pd.read_csv("Car_Price_dataset.csv")
dataset

```

	car_ID	symboling	CarName	fueltype	
aspiration \					
0	1	3	alfa-romero giulia	gas	std
1	2	3	alfa-romero stelvio	gas	std
2	3	1	alfa-romero Quadrifoglio	gas	std
3	4	2	audi 100 ls	gas	std
4	5	2	audi 100ls	gas	std
..
200	201	-1	volvo 145e (sw)	gas	std
201	202	-1	volvo 144ea	gas	turbo
202	203	-1	volvo 244dl	gas	std
203	204	-1	volvo 246	diesel	turbo
204	205	-1	volvo 264gl	gas	turbo
doornumber	carbody	drivewheel	engine	location	wheelbase ...
\					
0	two	convertible	rwd	front	88.6 ...
1	two	convertible	rwd	front	88.6 ...
2	two	hatchback	rwd	front	94.5 ...
3	four	sedan	fwd	front	99.8 ...
4	four	sedan	4wd	front	99.4 ...

..
200	four	sedan	rwd	front	109.1	...
201	four	sedan	rwd	front	109.1	...
202	four	sedan	rwd	front	109.1	...
203	four	sedan	rwd	front	109.1	...
204	four	sedan	rwd	front	109.1	...

	engine	size	fuel	system	boreratio	stroke	compressionratio
horsepower \							
0	130	mpfi	3.47	2.68	9.0		
111							
1	130	mpfi	3.47	2.68	9.0		
111							
2	152	mpfi	2.68	3.47	9.0		
154							
3	109	mpfi	3.19	3.40	10.0		
102							
4	136	mpfi	3.19	3.40	8.0		
115							
..
...							
200	141	mpfi	3.78	3.15	9.5		
114							
201	141	mpfi	3.78	3.15	8.7		
160							
202	173	mpfi	3.58	2.87	8.8		
134							
203	145	idi	3.01	3.40	23.0		
106							
204	141	mpfi	3.78	3.15	9.5		
114							

	peakrpm	citympg	highwaympg	price
0	5000	21.0	27.0	13495.0
1	5000	21.0	27.0	16500.0
2	5000	NaN	26.0	16500.0
3	5500	24.0	30.0	13950.0
4	5500	18.0	22.0	17450.0
..
200	5400	23.0	28.0	16845.0
201	5300	19.0	25.0	19045.0
202	5500	18.0	23.0	21485.0
203	4800	26.0	27.0	22470.0

```
204      5400      19.0      25.0  22625.0
```

```
[205 rows x 26 columns]
```

```
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer
```

```
text_features = ["CarName", "doornumber", "cylindernumber"]
oneHot = OneHotEncoder()
```

```
transformer = ColumnTransformer([('oneHot', oneHot, text_features)])
X = transformer.fit_transform(dataset)
y = dataset['price']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.25)
reg = LinearRegression()
reg.fit(X_train, y_train)
y_preds = reg.predict(X_test)
```

```
mae = mean_absolute_error(y_test, y_preds)
r2 = r2_score(y_test, y_preds)
mse = mean_squared_error(y_test, y_preds)
```

```
mae, r2, mse
#(2755.219116212158, 0.7459019410391048, 15174421.15204328)
#(2357.083638592613, 0.8334032631319556, 11006688.483041983)
(4615.984445080328, 0.45784338089448595, 46032922.96933926)
```

```
dataset.dtypes
```

car_ID	int64
symboling	int64
CarName	object
fueltype	object
aspiration	object
doornumber	object
carbody	object
drivewheel	object
enginelocation	object
wheelbase	float64
carlength	float64
carwidth	float64
carheight	float64
curbweight	float64
enginetype	object
cylindernumber	object
enginesize	int64
fuelsystem	object
bore ratio	float64

